

Deploying and Managing Microservices in a Cloud-Native Environment

Objective:

This project aims to migrate a monolithic application to a Kubernetes-based microservices architecture to achieve better scalability, high availability, and employment flexibility. As a DevOps Engineer at ABVC Solutions, the focus is on setting up Kubernetes infrastructure, containerizing services, and implementing cloud-native deployment practices.

Infrastructure and Tools Used:

- Kubernetes Cluster: Provisioned using EKS (Elastic Kubernetes Service)
- Containerization: Docker (with images stored on Docker Hub)
- CI/CD & CLI Tools: kubectl, docker, hey (load generation), curl
- Monitoring/Scaling: Horizontal Pod Autoscaler (HPA)
- Storage: Kubernetes Persistent Volumes (PV) and Persistent Volume Claims (PVC)

Implementation Steps:

1. Kubernetes Cluster Setup

- Created Kubernetes cluster using EKS on AWS
- Installed and configured kubectl to interact with the cluster:
`aws eks --region <region> update-kubeconfig --name <cluster-name>`
`kubectl get nodes`

2. Containerization & Docker

- Built a sample microservices application with three services:
 1. User Service
 2. Product Service
 3. Order Service
- Created separate Dockerfiles for each service
- Built and pushed images to Docker Hub:
`docker build -t <username>/user-service:latest .`
`docker push <username>/user-service:latest`

3. Kubernetes Deployments

- Created Deployment YAML files for each service with:

1. Container image

2. Replicas

3. Ports

4. Env variables

- Deployed with:

kubectl apply -f deployments/

4. Service Discovery & Load Balancing

- Defined Kubernetes Service resources:

1. Used ClusterIP for internal microservice communication

2. Used NodePort for external access during development

- Verified communication using internal DNS and curl requests

5. Scaling with HPA

- Enabled autoscaling for Product Service using HPA:

apiVersion: autoscaling/v2

kind: HorizontalPodAutoscaler

...

- Simulated load with:

hey -z 1m -c 50 http://<node-ip>:<port>/products

- Observed scaling activity with:

kubectl get hpa

6. Persistent Storage

- Created Persistent Volume (PV) and Persistent Volume Claim (PVC)
- Configured Order Service to store and read data from mounted volume
- Verified that data persisted even after pod restarts

Key Points :

1. Cluster Setup

- Used EKS for managed Kubernetes
- Configured access with AWS CLI and kubectl

2. Deployment & Testing

- Dockerized and deployed 3 services
- Set up internal/external service exposure
- Verified inter-service communication

3. Tools & Technologies

- AWS EKS, Docker, Kubernetes, Helm (if used), HPA, PV/PVC, kubectl, curl, hey

4. Challenges Faced

Challenge	Solution Implemented
Docker image size optimization	Used multi-stage builds in Dockerfiles
Inter-service DNS issues	Ensured proper service names and namespace usage
Autoscaler not triggering	Increased CPU load using stress/load tools
Data loss after restart	Implemented PVC and volume mounts correctly in yaml

Outcome:

The project successfully demonstrated:

- End-to-end deployment of containerized microservices
- Kubernetes concepts such as services, scaling, and persistent storage
- Real-world microservice orchestration and fault tolerance

```
ubuntu@ip-172-31-88-41: ~/microservices/product-service
-cpus          Number of used cpu cores.
               (default for current machine is 1 cores)
ubuntu@ip-172-31-88-41:~/microservices/product-service$ hey -z 2m -c 10 http://3.94.80.138:5000/

Summary:
Total:      120.0054 secs
Slowest:    0.0000 secs
Fastest:    0.0000 secs
Average:    NaN secs
Requests/sec: 0.5000

Response time histogram:

Latency distribution:

Details (average, fastest, slowest):
DNS+diatup:  NaN secs, 0.0000 secs, 0.0000 secs
DNS-lookup:  NaN secs, 0.0000 secs, 0.0000 secs
req write:   NaN secs, 0.0000 secs, 0.0000 secs
resp wait:   NaN secs, 0.0000 secs, 0.0000 secs
resp read:   NaN secs, 0.0000 secs, 0.0000 secs

Status code distribution:

Error distribution:
[46] Get "http://3.94.80.138:5000/": context deadline exceeded (Client.Timeout exceeded while awaiting headers)
[14] Get "http://3.94.80.138:5000/": dial tcp 3.94.80.138:5000: i/o timeout (Client.Timeout exceeded while awaiting headers)

ubuntu@ip-172-31-88-41:~/microservices/product-service$ kubectl get hpa
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
product-service     Deployment/product-service  cpu: <unknown>/50%  1         5         2          10m
ubuntu@ip-172-31-88-41:~/microservices/product-service$
```

```
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
order-service       1/1     1             1           37m
product-service     0/2     2             0           108m
user-service        0/2     2             0           127m
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
order-service       1/1     1             1           37m
product-service     0/2     2             0           108m
user-service        0/2     2             0           127m
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get svc
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)
kubernetes          ClusterIP           10.100.0.1      <none>            443/TCP
order-service       LoadBalancer       10.100.88.168   ac9f6cb3195ae479c98df18d1a42ce38-1996437556.us-east-1.elb.amazonaws.com 5000:31451/TCP
product-service     LoadBalancer       10.100.24.80    a6b9f0c53dfbd4f29930fcc5cba2b883-14814814.us-east-1.elb.amazonaws.com 5000:32449/TCP
user-service        ClusterIP           10.100.179.158  <none>            5000/TCP
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get hpa
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
product-service     Deployment/product-service  cpu: <unknown>/50%  1         5         2          108m
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get pv
NAME                CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM                STORAGECLASS   VOLUMEATTRIBUTESCLASS   REASON   AGE
order-pv            1Gi        RWO            Retain           Bound    default/order-pvc    <unset>        <unset>                  87m
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get pvc
NAME                STATUS     VOLUME          CAPACITY   ACCESS MODES   STORAGECLASS   VOLUMEATTRIBUTESCLASS   AGE
order-pvc           Bound     order-pv        1Gi        RWO            <unset>        <unset>                  86m
ubuntu@ip-172-31-88-41:~/microservices$ kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
```

```
ubuntu@ip-172-31-88-41: ~/microservices/k8s/order-service
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ cd ..
ubuntu@ip-172-31-88-41:~/microservices/k8s$ cd..
cd..: command not found
ubuntu@ip-172-31-88-41:~/microservices/k8s$ cd ..
ubuntu@ip-172-31-88-41:~/microservices$ cd product-service/
ubuntu@ip-172-31-88-41:~/microservices/product-service$ ls
Dockerfile  app.py  product-deployment.yaml  product-service.yaml
ubuntu@ip-172-31-88-41:~/microservices/product-service$ cd ../user-service/
ubuntu@ip-172-31-88-41:~/microservices/user-service$ ls
Dockerfile  app.py  requirements.txt  user-deployment.yaml  user-service.yaml
ubuntu@ip-172-31-88-41:~/microservices/user-service$ cd ..
ubuntu@ip-172-31-88-41:~/microservices$ ls
k8s  nano.8830.save  product-service  user-service
ubuntu@ip-172-31-88-41:~/microservices$ cd k8s/
ubuntu@ip-172-31-88-41:~/microservices/k8s$ cd order-service/
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ nano order-service.yaml
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ kubectl apply -f order-service.yaml
error: the path "k8s/order-service/order-deployment.yaml" does not exist
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ kubectl apply -f order-deployment.yaml
deployment.apps/order-service unchanged
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ kubectl apply -f k8s/order-service/order-service.yaml
error: the path "k8s/order-service/order-service.yaml" does not exist
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ kubectl apply -f order-service.yaml
service/order-service created
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$ kubectl get svc order-service
NAME          TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)
order-service  LoadBalancer  10.100.88.168    ac9f6cb3195ae479c98df18d1a42ce38-1996437556.us-east-1.elb.amazonaws.com  5000:31451/TC
P 14s
ubuntu@ip-172-31-88-41:~/microservices/k8s/order-service$
```

